



LINUX
SECURITY
SUMMIT
EUROPE

Using the ARM Trust-Zone to control Tamper Resistant Processors

U-boot/Linux Cryptographic Support

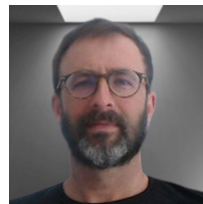
Jorge Ramirez-Ortiz
Device Security Lead

@Foundries.io



AGENDA

1. ARM TrustZone
2. Runtime Resident Firmware
3. OP-TEE: Root of Trust
4. OP-TEE: Secure Persistent Storage
5. OP-TEE: Trusted Application
6. OP-TEE: PKCS#11 Trusted Application
7. OP-TEE: Executing in the TrustZone
8. OP-TEE: Trusted Thread Exit
9. Yielding SMC: Working in the TrustZone
10. Use Case: NXP SE050X HSM (I2C)
11. Use Case: Versal ACAP Crypto
12. References

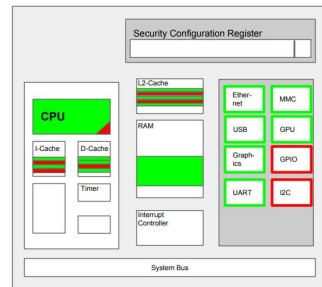


Maintainership in OP-TEE:
Versal ACAP and drivers
ZynqMP drivers
SE05X crypto driver
iMX I2C driver
iMX RNG driver

jorge@foundries.io

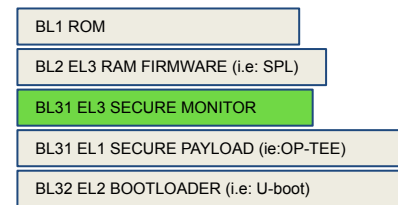
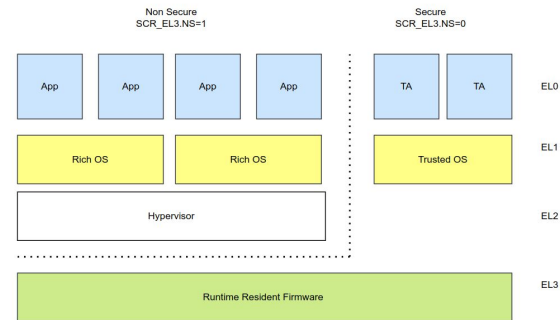
ARM TrustZone

- Hardware feature on Cortex-A and since recently Cortex-M profiles
- Hardware enforced isolation built into the CPU
 - Additional Processor Mode (CPR15 NS bit)
 - MMU mode awareness (cache/TLB entries can coexist)
 - New Exception Level (EL3) to handle mode transition
 - Runtime Resident Firmware
 - Runtime Resident Firmware Services (std_svc or spd)
 - Vendor Specific Peripheral partitioning (i.e GPIO)
- TrustZone must be handled with care
 - ARM Trusted Firmware project
- QEMU emulation has been supported for nearly 10 years now




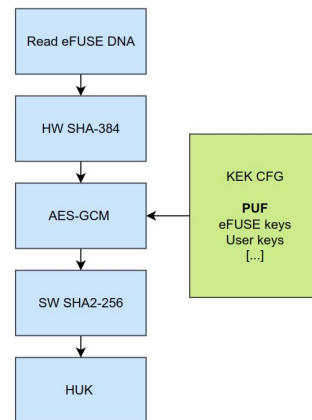
Runtime Resident Firmware

- Runtime Resident Firmware - executes at EL3
- Should be installed under the Root of Trust
- Entered via the SMC privileged call
 - smccc - follows calling convention
 - different calls:
 - Fast: Atomic - ARM, CPU, SiP(silicon), OEM ..
 - **Yielding**: Can be suspended.
- OP-TEE Runtime Resident Firmware
- TrustedFirmware.org Runtime Resident Firmware
 - Provides Runtime Services:
 - Secure Payload Dispatcher
 - » switches between normal and secure world
 - » handles SMC functions targeting S-EL1
 - » manages S-EL1 context



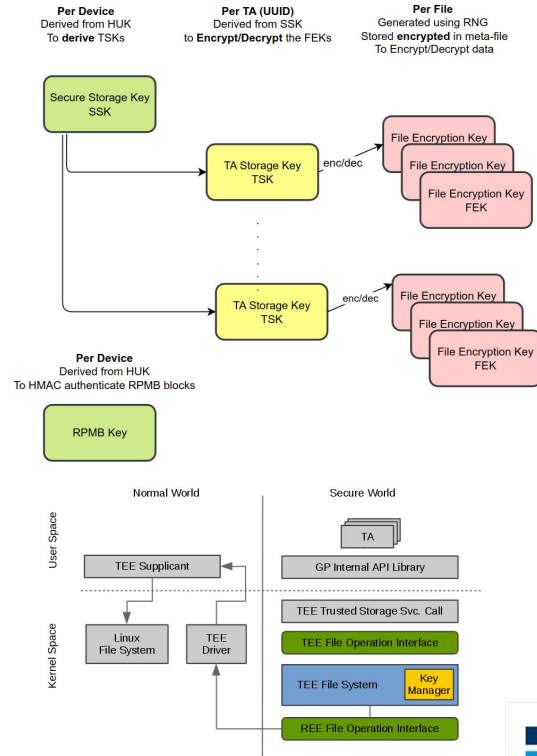
OP-TEE - Root of Trust

- **Hardware Unique Key** 128 bits
 - per device identifier
 - symmetric key used to derive (KDF) all other keys
 - secure storage keys
 - securing the linux keyring: sealing keys
 - external HSM protocol encryption keys (NXP SE05x)
 - must not be readable outside the TrustZone
 - risks secure storage duplication and its decryption
- Could be
 - Provided by SoC (iMX CAAM)
 - Generated by OP-TEE driver (AMD Versal ACAP) 
 - uid + AES-GCM
- Serves as a device unique identifier
- Must only be available when booting signed firmware



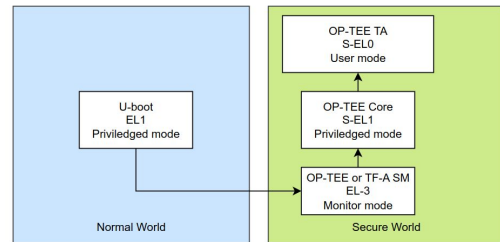
OP-TEE - Secure Persistent Storage

- Global Platform Trusted Storage Requirements
 - May be backed by NS resources
 - OP-TEE: REE-FS :
 - FEK **AES-GCM** auth enc/dec data
 - eMMC RPMB support
 - OP-TEE: RPMB-FS:
 - FEK **AES-CBC** enc/dec data
 - RPMB key: HMAC RPMB blocks
 - Bound to a device
 - Each TA can only access its own storage
 - Must protect against replay attacks

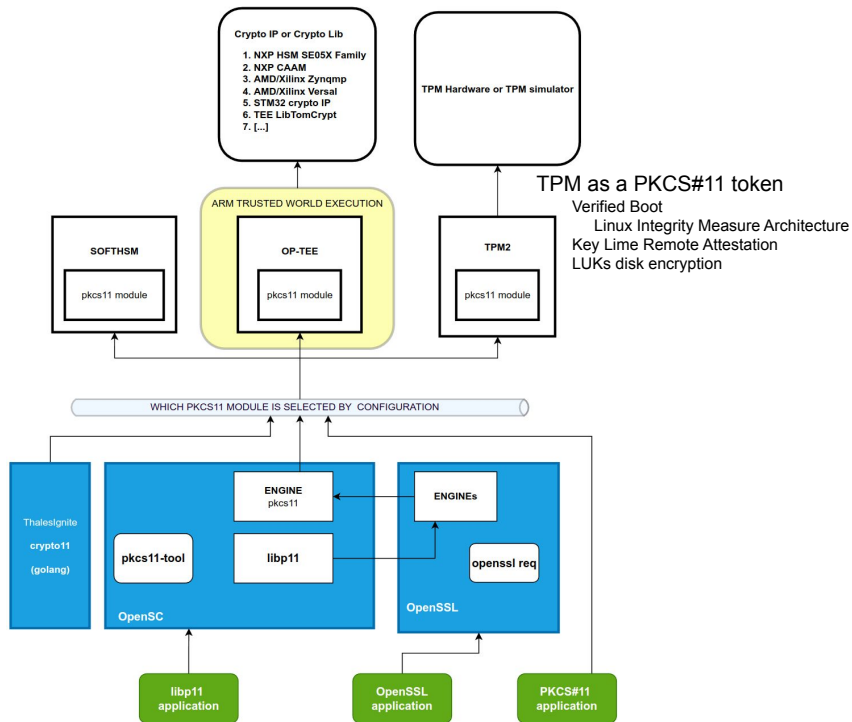


OP-TEE - Trusted Application

- The execution is split by the NS bit (TA is NS=0)
- TA are RSA signed
 - asymmetric public key built in the TEE core
- TA might be AES-GCM encrypted
 - symmetric key generation is platform dependent
- TA can be
 - stored in the REE filesystem (TA)
 - built-in the TEE image (Early TA)
- Upon load they are verified and decrypted
- TA are versioned (ta_ver.db) to prevent rollback attacks
 - However only effective when using RPMB_FS
- Built time TA configuration flags
 - multisession, keep alive, stack size, heap size ...



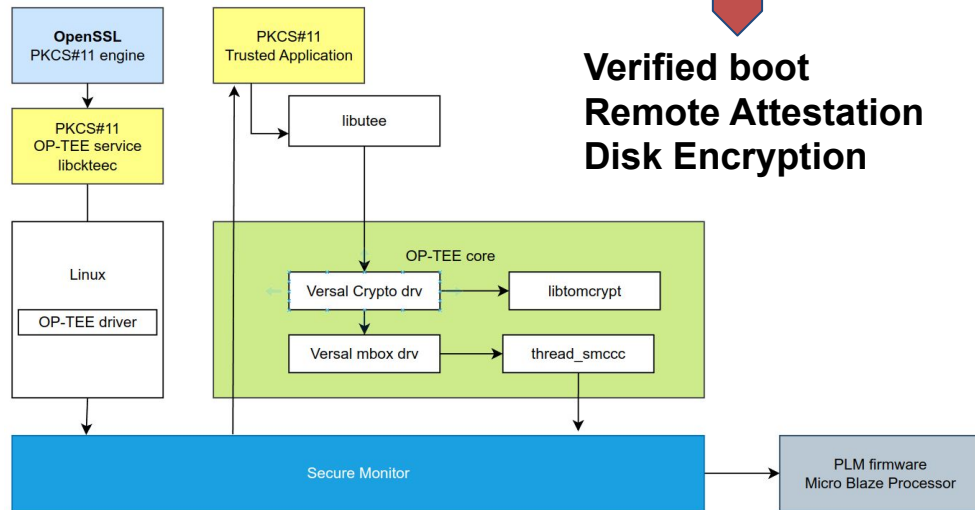
OP-TEE PKCS#11 Trusted Application



PKCS#11:

A generic secure interface for all things crypto
Key management, Crypto operations, Storage

Verified boot
Remote Attestation
Disk Encryption

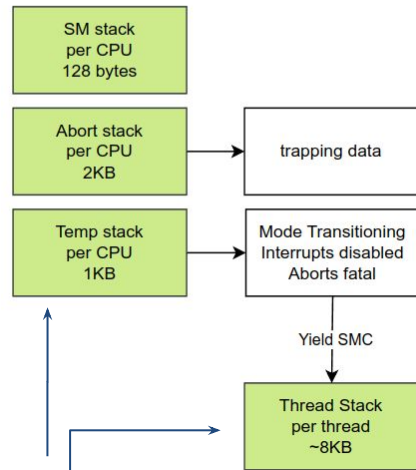


AGENDA

1. ARM TrustZone
2. Runtime Resident Firmware
3. OP-TEE: Root of Trust
4. OP-TEE: Secure Persistent Storage
5. OP-TEE: Trusted Application
6. OP-TEE: PKCS#11 Trusted Application
7. OP-TEE: Executing in the TrustZone
8. OP-TEE: Trusted Thread Exit
9. Yielding SMC: Working in the TrustZone
10. Use Case: NXP SE050X HSM (I2C)
11. Use Case: Versal ACAP Crypto
12. References

OP-TEE: Executing in the Trust Zone

- OP-TEE (optionally) implements an EL3 Runtime Resident Firmware (SM)
- A Secure Monitor manages all entries and exits of the secure world via SMC
- SMC
 - Fast: IRQ masked
 - Yielding: IRQ unmasked at some point
 - S-EL1 context can be suspended and resumed
 - TT scheduling controlled by REE
- OP-TEE does NOT implement a scheduler
 - Trusted Threads (shadow) are tied to the execution of the calling thread
 - All requests are always REE initiated
 - Number of TT is statically allocated
- OP-TEE supports multiple threads of execution
 - OP-TEE implements synchronization primitives
 - Native: spinlock - do not take with interrupts enabled
 - REE: mutex and condition variables
 - Do not take in interrupt context, do not leave the TA holding a mutex
 - In REE Linux it uses the Linux *completion* structures



Size configurable at build time
(i.e: NXP SE05x requires 16KB)

OP-TEE: Trusted Thread exit

- Every execution request initiates in the Normal World
- Yielding SMC calls context can be restored → Sync/Async TT exits
- Asynchronous Exits
 - Foreign interrupts will suspend TT and return to NS
- Synchronous Exits
 - Callback the monitor software at EL3
 - `thread_smccc`
 - Callback the REE at EL1 via monitor software at EL3
 - `thread_rpc(OPTEE_RPC_CMD ...)`
 - the REE can handle the request at EL1
 - U-boot can access RPMB
 - Linux at EL-1 can provide I2C access to the TEE
 - or notify an EL0 daemon to handle the request
 - the EL0 daemon is known as the *tee_suppllicant*
 - » *tee_suppllicant* handles RPMB requests in Linux



Normal Exit - `tee_entry_std()` returns

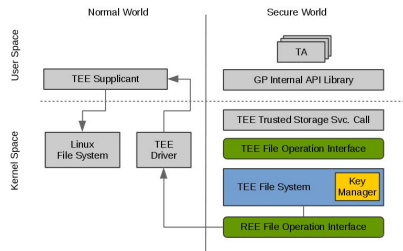
Foreign Interrupt Exit
Interrupt targeting Normal World
Interrupt targeting Secure Monitor

Thread RPC Exit

Thread SMC Exit

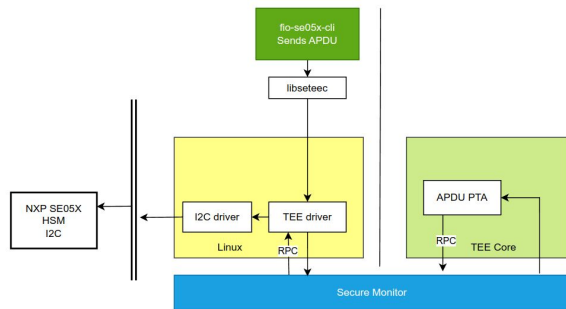
Yielding SMC: Working in the TrustZone

- EL0: User Linux, TEE supplicant



Example: Write to eMMC RPMB

- EL0: User Linux, NO TEE supplicant



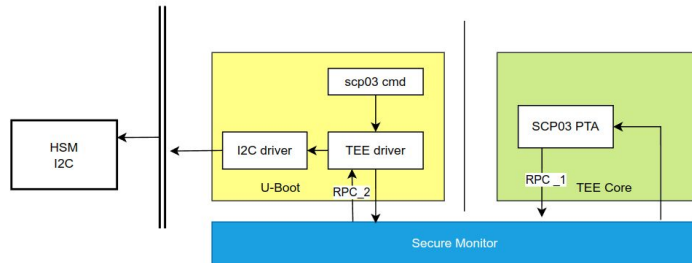
Example: Access to Secure Element

<https://github.com/foundriesio/fio-se05x-cli>

drivers/tee/optee/rpc.c

Not all RPCs on Linux need to have support from the TEE supplicant

- EL1: Single threaded Bootloader, U-boot



Example: HSM Key Rotation, scp_03 cmd

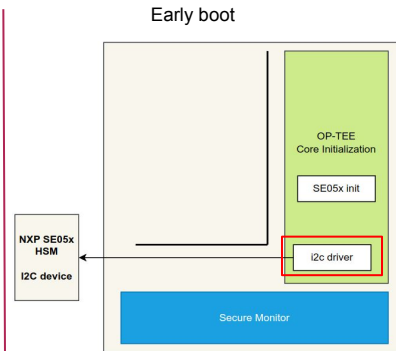
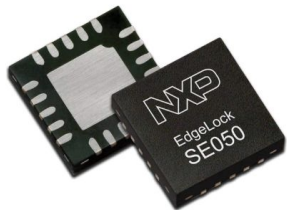
<https://u-boot.readthedocs.io/en/latest/usage/cmd/scp03.html>

SPL can load and boot OP-TEE so it is ready for U-boot

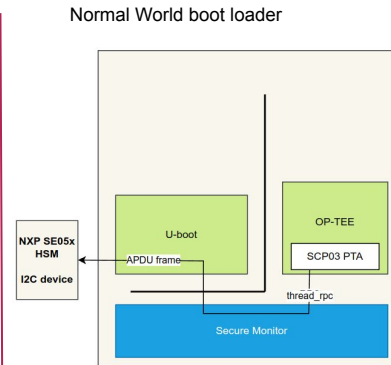
```
void optee_rpc_cmd(struct tee_context *ctx, struct optee *optee,
                  struct optee_arg *arg)
{
    switch (arg->cmd) {
        case OPTEE_RPC_CMD_GET_TIME:
            handle_rpc_func_cmd_get_time(arg);
            break;
        case OPTEE_RPC_CMD_NOTIFICATION:
            handle_rpc_func_cmd_notification(arg);
            break;
        case OPTEE_RPC_CMD_SUSPENDING:
            handle_rpc_func_cmd_suspend(arg);
            break;
        case OPTEE_RPC_CMD_I2C_TRANSFER:
            handle_rpc_func_cmd_i2c_transfer(ctx, arg);
            break;
        default:
            handle_rpc_supp_cmd(ctx, optee, arg);
    }
}
```

Use Case: NXP SE05X HSM (I2C)

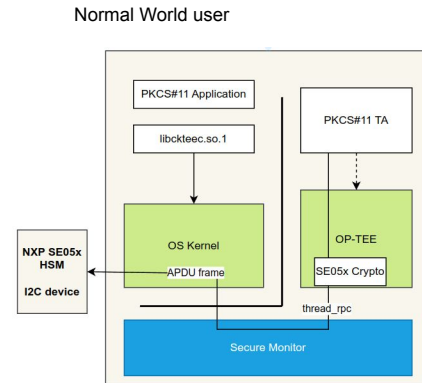
- THREAD_RPC: delegate to EL1
- NXP SE050x Root of Trust at IC level
- FIPS 140-2 Certified Security device
- Keys can be pre-provisioned in memory
- Pre-provisioned keys can be imported and handled by OP-TEE's PKCS#11
- Provisioning (SCP03)
 - Data on i2c bus is **AES-GCM** with session keys.
 - AES-GCM in SW (LibTomCrypt)
 - Keys derived from programmable NVM values
 - KDF uses the HUK
 - HUK must be stable



Device Provisioning
Unique identifier
GP SCP03 enablement
Key rotation



Device Provisioning
Might update SCP03 keys



Secure Storage
Crypto Operations
RSA
ECC
TRNG
AES-CTR



Controlling access to the device during the boot sequence

```
int glue_i2c_init(void)
{
    if (transfer == xpc_i2c_transfer)
        return 0;

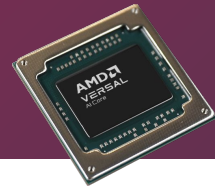
    transfer = native_i2c_transfer;
    return native_i2c_init();
}

static TEE_Result load_trampoline(void)
{
    if (IS_ENABLED(CONFIG_CORE_SE05X_I2C_TRAMPOLINE))
        transfer = xpc_i2c_transfer;

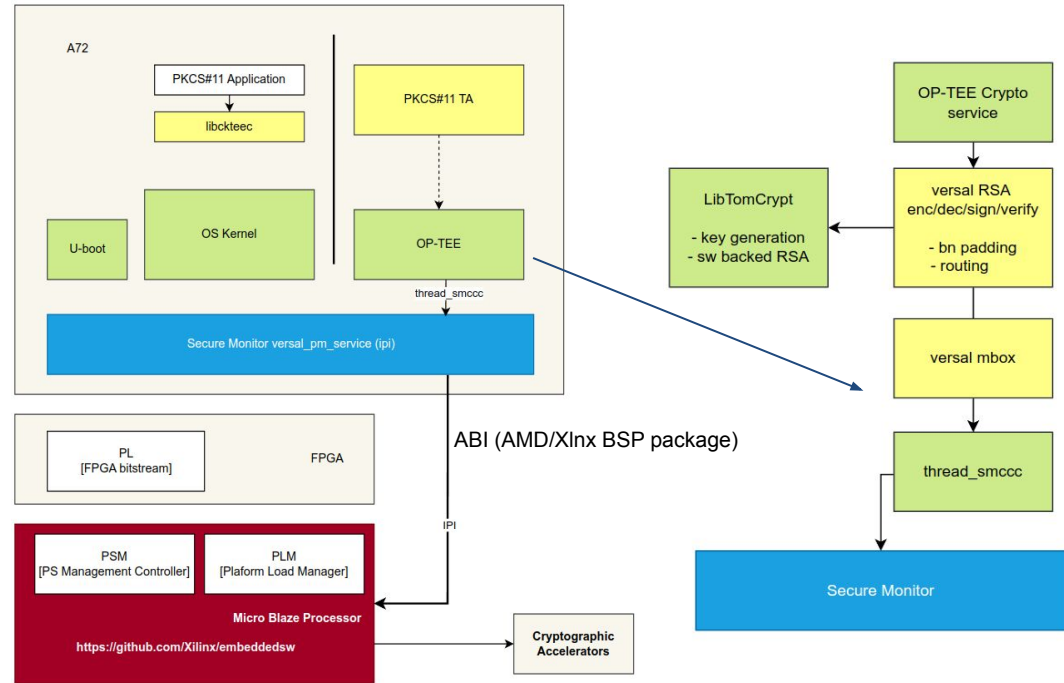
    return TEE_SUCCESS;
}

boot_final(load_trampoline);
```

Use Case: Versal ACAP Crypto



- THREAD_SMCCC delegate to EL3
- Most Secure Services are implemented in the **PLM firmware** executing in a remote processor
 - Cryptography
 - AES-GCM
 - RSA/ECC
 - SHA3-384
 - Physical Unclonable Device
 - eFuse
 - FPGA configuration
- Others are implemented natively in OP-TEE
 - TRNG, GPIO
- The Mailbox driver is implemented in the **Secure Monitor** firmware routing to the PLM
 - Defines ABI with PLM services
- Device tree: critical services in use by the Secure World must be reserved and its **secure-status** property set accordingly.[1]



[1] <https://www.kernel.org/doc/Documentation/devicetree/bindings/arm/secure.txt>

References

- SMC Calling Convention : <https://developer.arm.com/documentation/den0028/latest>
- OP-TEE: <https://optee.readthedocs.io/en/latest/>
 - QEMU ARMv8 <https://optee.readthedocs.io/en/latest/building/devices/qemu.html#qemu-v8>
- AMD Versal ACAP : <https://optee.readthedocs.io/en/latest/building/devices/versal.html>
 - Crypto Driver: https://github.com/OP-TEE/optee_os/tree/master/core/drivers/crypto/versal
 - **AMD Versal ACAP, Technical Overview, Embedded World 2023**, <https://shorturl.at/gLOQ0>
- SE05X: <https://optee.readthedocs.io/en/latest/architecture/crypto.html#nxp-se05x-family-of-secure-elements>
 - Crypto Driver: https://github.com/OP-TEE/optee_os/tree/master/core/drivers/crypto/se050
 - **Secure Elements in a TEE, Embedded Recipes 2022**, <https://shorturl.at/txFW3>



LINUX SECURITY SUMMIT

EUROPE

Secure Boot on TEE

SoC fuses/OTP registers store the Root of Trust

- The public key hash to save register space.
- The first bootloader MUST be signed with the RoT key
 - The public key is then embedded in the image for ROM verification

The first bootloader then boots a FIT image [multiple signed binaries in a single file]

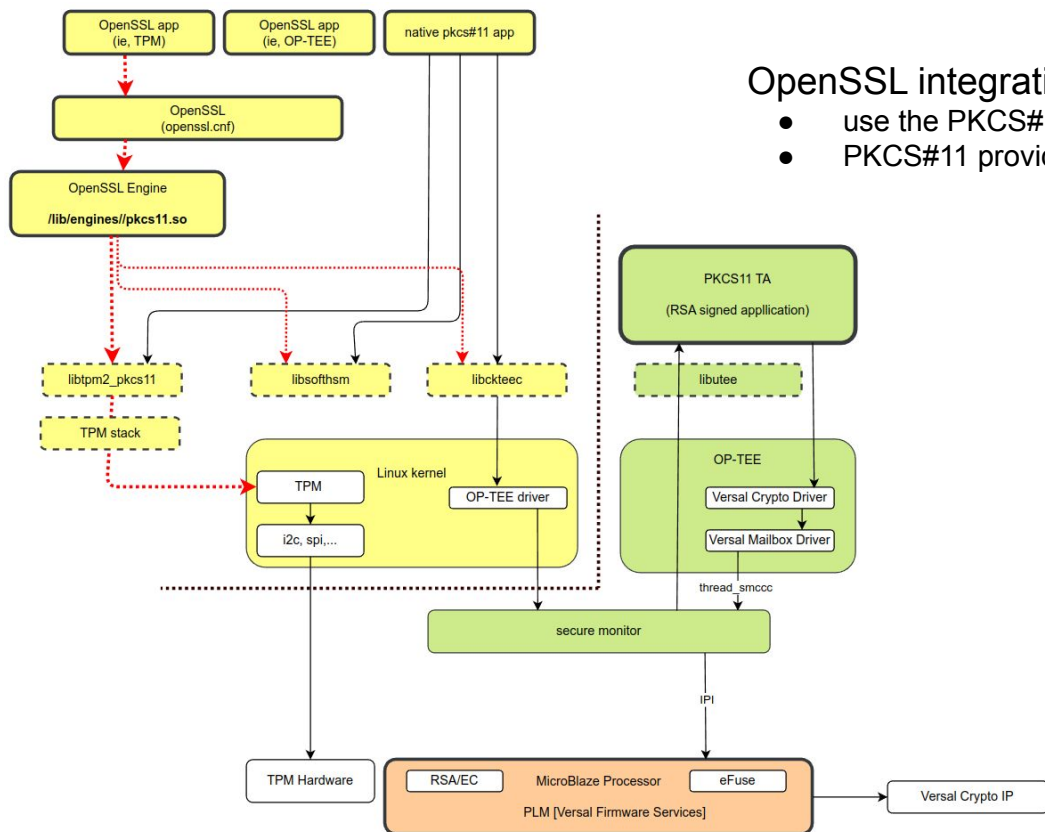
- Images in the FIT should be **signed with different keys** to reduce the attack surface **and enable key rotation**
 - Reduce RoT key exposure [unless the FIT is also being validated by the ROM]
- The First bootloader will validate and boot the subsequent firmwares.
 - Usually a software implementation of the cryptographic driver.

Once the TEE is up early in the boot process (**unless you are a ChromeOS user**) it can access the IP to perform cryptographic operations. For example

- AES-GCM (for secure storage)
- RSA verification of the TAs

U-Boot can then access OP-TEE for further validation and Secure Storage in RPMB.

BACKUP



OpenSSL integration

- use the PKCS#11 engine
- PKCS#11 provider not quite ready

ARM64 Code dive

Normal World EL1

```
static void do_call_with_arg(struct udevice *dev, struct optee_msg_arg *arg)
{
    struct optee_pdata *pdata = dev_get_plat(dev);
    struct rpc_param param = { .a0 = OPTEE_SMC_CALL_WITH_ARG };
    void *page_list = NULL;

    reg_pair_from_64(&param.a1, &param.a2, virt_to_phys(arg));
    while (true) {
        struct arm_smccc_res res;

        /* If cache are off from U-Boot, sync the cache shared with OP-TEE */
        if (l1dcache_status())
            flush_shm_dcache(dev, arg);

        pdata->invoke_fn(param.a0, param.a1, param.a2, param.a3,
                        param.a4, param.a5, param.a6, param.a7, &res);

        /* If cache are off from U-Boot, sync the cache shared with OP-TEE */
        if (l1dcache_status())
            flush_shm_dcache(dev, arg);

        free(page_list);
        page_list = NULL;

        if (OPTEE_SMC_RETURN_IS_RPC(res.a0)) {
            param.a0 = res.a0;
            param.a1 = res.a1;
            param.a2 = res.a2;
            param.a3 = res.a3;
            handle_rpc(dev, &param, &page_list);
        } else {
            /*
             * In case we've accessed RPMB to serve an RPC
             * request we need to restore the previously
             * selected partition as the caller may expect it
             * to remain unchanged.
             */
            optee_suppl_rpmb_release(dev);
            return call_err_to_res(res.a0);
        }
    }
}
```

Notifies the tee_supplcant (RPMB, I2C, ...)

Then calls the monitor again

Monitor EL3

opteed_smc_handler

```
/*
 * OPTEE is returning from a call or being preempted from a call, in
 * either case execution should <u>continue in the normal world</u>.
 */
void tee_smc_opteed_return_call_done()
{
    /* This is the result from the secure client of an
     * earlier request. The results are in x0-x2. Copy it
     * into the non-secure context, save the secure state
     * and return to the non-secure state.
     */
    assert(handle == cn_get_context(SECURE));
    cn_all_sysregs_context_save(SECURE);

    /* Get a reference to the non-secure context */
    ns_cpu_context = cn_get_context(NON_SECURE);
    assert(ns_cpu_context);

    /* Restore non-secure state */
    cn_all_sysregs_context_restore(NON_SECURE);
    cn_set_next_err_context(NON_SECURE);

    SMC_RET0(ns_cpu_context, x1, x2, x3, x0);
}
```

```
/* void thread_smccc(struct thread_smc_args *arg_res) */
FUNC thread_smccc, 1
push    x0, xzr
mov     x0, x0
load_xregs x0, 0, 0, 7
#ifdef CFG_CORE_SECURE_SMC
hvc     #0
#else
smc     #0
#endif
pop     x0, xzr
store_xregs x0, 0, 0, 7
ret

END_FUNC thread_smccc
```

```
/* void thread_rpc_spr(uint32_t rv[THREAD_RPC_NUM_ARGS], uint64_t spsr) */
FUNC thread_rpc_spr, 1
/* Mask all maskable exceptions before switching to temporary stack */
msr     daifset, #DAIFBIT_ALL
push    x0, xzr
push    x1, x30
bl      thread_get_ctx_regs
ldr     x30, [sp, #0]
store_xregs x0, THREAD_CTX_REGS_X19, 19, 30
mov     x19, x0

#ifdef CFG_CORE_PAUTH
/* Save APIAKEY */
read_apikeyh1 x1
read_apikeylo x2
store_xregs x0, THREAD_CTX_REGS_APIAKEY_HI, 1, 2
#endif

#endif

bl      thread_get_tmp_sp
pop     x1, xzr /* Match "push x1, x30" above */
mov     x2, sp
str     x2, [x19, #THREAD_CTX_REGS_SP]
ldr     x20, [sp] /* Get pointer to rv[] */
mov     sp, x0 /* Switch to tmp stack */

/*
 * We need to read rv[] early, because thread_state_suspend
 * can invoke virt_unset_guest() which will unmap pages,
 * & where rv[] resides
 */
load_wregs x20, 0, 21, 23 /* Load rv[] into w20-w22 */

adr     x2, thread_rpc_return
mov     w0, #THREAD_FLAGS_COPY_ARGS_ON_RETURN
bl      thread_state_suspend
mov     x4, x0 /* Supply thread index */
ldr     w0, [TEESMC_OPTEE_RETURN_CALL_DONE]
mov     x1, x21
mov     x2, x22
mov     x3, x23
smc     #0

/* SMC should not return */
panic_at_smc_return
```

S-EL1

```
int thread_state_suspend(uint32_t flags, uint32_t cpsr, vaddr_t pc)
{
    struct thread_core local *l = thread_get_core_local();
    int ct = l->curr_thread;
    assert(ct != THREAD_ID_INVALID);

    if (core_mmu_user_mapping_is_active())
        ftrace_suspend();

    thread_check_cannaries();
    release_unused_kernel_stack(threads + ct, cpsr);

    if (is_from_user(cpsr)) {
        thread_user_save_vfp();
        tee_ta_update_session_utm_suspend();
        tee_ta_gprof_sample_pc(pc);
    }
    thread_lazy_restore_ns_vfp();

    thread_lock_global();

    assert(threads[ct].state == THREAD_STATE_ACTIVE);
    threads[ct].flags |= flags;
    threads[ct].regs.cpsr = cpsr;
    threads[ct].regs.pc = pc;
    threads[ct].state = THREAD_STATE_SUSPENDED;

    threads[ct].have_user_map = core_mmu_user_mapping_is_active();
    if (threads[ct].have_user_map) {
        if (threads[ct].flags & THREAD_FLAGS_EXIT_ON_FOREIGN_INTR)
            tee_ta_ftrace_update_times_suspend();
        core_mmu_get_user_map(&threads[ct].user_map);
        core_mmu_set_user_map(NULL);
    }

    if (IS_ENABLED(CFG_SECURE_PARTITION)) {
        struct ts_session *ts_sess =
            TAILQ_FIRST(&threads[ct].tsd.sess_stack);

        smpc_sp_set_to_preempted(ts_sess);
    }

    l->curr_thread = THREAD_ID_INVALID;

    if (IS_ENABLED(CFG_NS_VIRTUALIZATION))
        virt_unset_guest();

    thread_unlock_global();

    return ct;
}
```

Zynqmp: U-boot/Linux

